

VPS Customer churn prediction - Part 2

JG Pardyak

02-9-2021

Introduction

Outlier detection

Tune `decision_tree` hyperparameters

Tune `rand_forest` hyperparameters

Model deployment

Further steps

Introduction

Motivation

This is the continuation of the presentations:


- ▶ <https://github.com/JacekPardyak/vps/blob/master/vps-part-1.pdf> ,
- ▶ <https://github.com/JacekPardyak/vps/blob/master/vps-part-2.pdf> .

In this presentation we check:


- ▶ outliers - data points that differs significantly from other observations,
- ▶ Feature engineering - combining and transforming further existing variables,
- ▶ Model tuning of `decision_tree` and `rand_forest`

Outlier detection


disk_ops variables




disk_octets variables



cpu variables



network variables



Tune decision_tree hyperparameters

Split and prepare data

```
set.seed(123)
vps_split <- initial_split(vps, strata = is_churn)
vps_train <- training(vps_split)
vps_test  <- testing(vps_split)
vps_recipe <- vps_train %>%
  recipe(is_churn ~ .) %>%
  step_rm(id) %>%
  step_corr(all_predictors()) #>%
  # make vars to be of mean zero
# step_center(all_predictors(), -all_outcomes()) %>%
# make vars to be standard dev of 1
#step_scale(all_predictors(), -all_outcomes())
```

Create model specification for tuning

```
tune_spec <-  
  decision_tree(  
    cost_complexity = tune(),  
    tree_depth = tune()  
  ) %>%  
  set_engine("rpart") %>%  
  set_mode("classification")
```

```
tune_spec
```

```
## Decision Tree Model Specification (classification)  
##  
## Main Arguments:  
##   cost_complexity = tune()  
##   tree_depth = tune()  
##  
## Computational engine: rpart
```

Create grid of hyperparameters values

```
tree_grid <- grid_regular(cost_complexity(),  
                           tree_depth(),  
                           levels = 5)  
  
tree_grid
```

```
## # A tibble: 25 x 2  
##   cost_complexity tree_depth  
##           <dbl>      <int>  
## 1  0.0000000001          1  
## 2  0.0000000178          1  
## 3  0.00000316            1  
## 4  0.000562              1  
## 5  0.1                   1  
## 6  0.0000000001          4  
## 7  0.0000000178          4  
## 8  0.00000316            4  
## 9  0.000562              4  
## 10 0.1                   4  
## # with 15 more rows
```

Create cross-validation folds for tuning

```
set.seed(234)
vps_folds <- vfold_cv(vps_train)
vps_folds
```

```
## # 10-fold cross-validation
```

```
## # A tibble: 10 x 2
```

##	splits	id
##	<list>	<chr>
## 1	<split [190/22]>	Fold01
## 2	<split [190/22]>	Fold02
## 3	<split [191/21]>	Fold03
## 4	<split [191/21]>	Fold04
## 5	<split [191/21]>	Fold05
## 6	<split [191/21]>	Fold06
## 7	<split [191/21]>	Fold07
## 8	<split [191/21]>	Fold08
## 9	<split [191/21]>	Fold09
## 10	<split [191/21]>	Fold10

Fit models at all the different values

```
set.seed(345)
tune_wf <- workflow() %>%
  add_model(tune_spec) %>%
  add_recipe(vps_recipe)
tree_res <-
  tune_wf %>%
  tune_grid(
    resamples = vps_folds,
    grid = tree_grid
  )

tree_res
```


```
## # Tuning results
## # 10-fold cross-validation
## # A tibble: 10 x 4
```

##	splits	id	.metrics	.notes
##	<list>	<chr>	<list>	<list>
##	1 <split [100/22]>	Fold01	<tibble [50 x 6]>	<tibble [50

Get metrics of different models

```
## # A tibble: 50 x 8
##   cost_complexity tree_depth .metric .estimator mean
##           <dbl>       <int> <chr>    <chr>    <dbl>
## 1  0.00000000001         1 accuracy binary    0.656
## 2  0.00000000001         1 roc_auc  binary    0.642
## 3  0.0000000178         1 accuracy binary    0.656
## 4  0.0000000178         1 roc_auc  binary    0.642
## 5  0.00000316           1 accuracy binary    0.656
## 6  0.00000316           1 roc_auc  binary    0.642
## 7  0.000562             1 accuracy binary    0.656
## 8  0.000562             1 roc_auc  binary    0.642
## 9  0.1                   1 accuracy binary    0.656
## 10 0.1                   1 roc_auc  binary    0.642
## # ... with 40 more rows
```


Plot metrics of different models



Show models with best metrics

```
## # A tibble: 5 x 8
##   cost_complexity tree_depth .metric .estimator mean
##           <dbl>       <int> <chr>    <chr>    <dbl>
## 1  0.00000000001         1 accuracy binary    0.656
## 2  0.0000000178         1 accuracy binary    0.656
## 3  0.00000316           1 accuracy binary    0.656
## 4  0.000562             1 accuracy binary    0.656
## 5  0.1                  1 accuracy binary    0.656
```

Pick one model with the best metrics

```
best_tree <- tree_res %>%  
  select_best("accuracy")  
best_tree
```

```
## # A tibble: 1 x 3  
##   cost_complexity tree_depth .config  
##           <dbl>       <int> <chr>  
## 1      0.0000000001           1 Preprocessor1_Model01
```

Finalizing best model

```
final_wf <-  
  tune_wf %>%  
  finalize_workflow(best_tree)
```

```
final_wf
```

```
## == Workflow =====  
## Preprocessor: Recipe  
## Model: decision_tree()  
##  
## -- Preprocessor -----  
## 2 Recipe Steps  
##  
## * step_rm()  
## * step_corr()  
##  
## -- Model -----  
## Decision Tree Model Specification (classification)  
##
```


The final fit

```
final_fit <-  
  final_wf %>%  
  last_fit(vps_split)
```

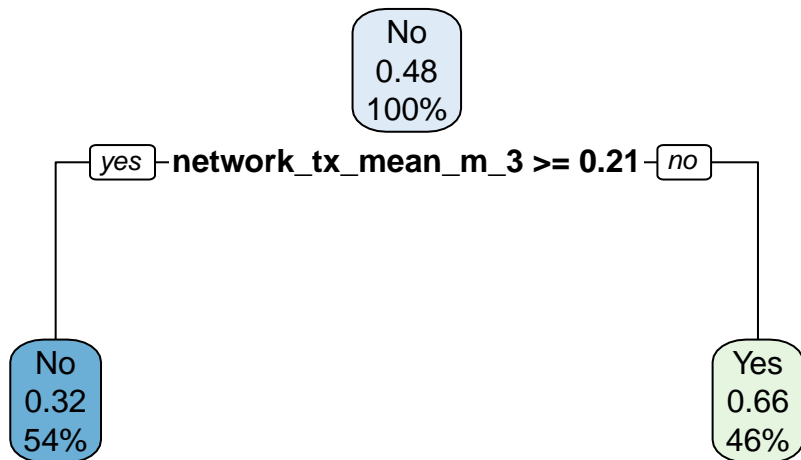
```
final_fit %>%  
  collect_metrics()
```

```
## # A tibble: 2 x 4  
##   .metric .estimator .estimate .config  
##   <chr>    <chr>          <dbl> <chr>  
## 1 accuracy binary          0.577 Preprocessor1_Model1  
## 2 roc_auc  binary          0.578 Preprocessor1_Model1
```


The final fit ROC curve



Visualizing decision tree from the workflow



Estimate variable importance



Tune `rand_forest` hyperparameters

Create model specification for tuning

```
tune_spec <- rand_forest(  
  mtry = tune(),  
  trees = 1000,  
  min_n = tune()  
) %>%  
  set_engine("ranger") %>%  
  set_mode("classification")
```

tune_spec

```
## Random Forest Model Specification (classification)  
##  
## Main Arguments:  
##   mtry = tune()  
##   trees = 1000  
##   min_n = tune()  
##  
## Computational engine: ranger
```

Create cross-validation folds for tuning

```
set.seed(234)
vps_folds <- vfold_cv(vps_train)
vps_folds
```

```
## # 10-fold cross-validation
```

```
## # A tibble: 10 x 2
```

##	splits	id
##	<list>	<chr>
##	1 <split [190/22]>	Fold01
##	2 <split [190/22]>	Fold02
##	3 <split [191/21]>	Fold03
##	4 <split [191/21]>	Fold04
##	5 <split [191/21]>	Fold05
##	6 <split [191/21]>	Fold06
##	7 <split [191/21]>	Fold07
##	8 <split [191/21]>	Fold08
##	9 <split [191/21]>	Fold09
##	10 <split [191/21]>	Fold10

Initially fit models at all the different values

```
set.seed(234)
doParallel::registerDoParallel()
tune_wf <- workflow() %>%
  add_recipe(vps_recipe) %>%
  add_model(tune_spec)

tune_res <- tune_grid(
  tune_wf,
  resamples = vps_folds,
  grid = 20
)
```


```
## i Creating pre-processing data to finalize unknown param
```

Get metrics of different models

```
## # A tibble: 40 x 8
```

```
##      mtry min_n .metric .estimator  mean      n std_err  
##    <int> <int> <chr>    <chr>    <dbl> <int>   <dbl>  
##  1      2      2 accuracy binary    0.571    10  0.0342  
##  2      2      2 roc_auc  binary    0.604    10  0.0399  
##  3     10     16 accuracy binary    0.562    10  0.0389  
##  4     10     16 roc_auc  binary    0.579    10  0.0435  
##  5      4     29 accuracy binary    0.599    10  0.0358  
##  6      4     29 roc_auc  binary    0.606    10  0.0426  
##  7      8     36 accuracy binary    0.585    10  0.0384  
##  8      8     36 roc_auc  binary    0.596    10  0.0404  
##  9      7     24 accuracy binary    0.566    10  0.0381  
## 10      7     24 roc_auc  binary    0.596    10  0.0417  
## # ... with 30 more rows
```

Plot initial tuning results



Create detailed grid of hyperparameters values

```
rf_grid <- grid_regular(  
  mtry(range = c(1, 3)),  
  min_n(range = c(1, 5)),  
  levels = 5  
)
```

```
rf_grid
```

```
## # A tibble: 15 x 2
```

```
##       mtry min_n
```

```
##   <int> <int>
```

```
## 1     1     1
```

```
## 2     2     1
```

```
## 3     3     1
```

```
## 4     1     2
```

```
## 5     2     2
```

```
## 6     3     2
```

```
## 7     1     3
```

```
## 8     2     3
```

Fit models at all the different values

```
set.seed(456)
forest_res <- tune_wf %>%
  tune_grid(
    resamples = vps_folds,
    grid = rf_grid
  )
```



Get metrics of different models

```
## # A tibble: 30 x 8
```


##	mtry	min_n	.metric	.estimator	mean	n	std_err		
##	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>	
##	1	1	1	accuracy	binary	0.576	10	0.0329	P
##	2	1	1	roc_auc	binary	0.612	10	0.0388	P
##	3	2	1	accuracy	binary	0.576	10	0.0297	P
##	4	2	1	roc_auc	binary	0.601	10	0.0426	P
##	5	3	1	accuracy	binary	0.566	10	0.0313	P
##	6	3	1	roc_auc	binary	0.588	10	0.0438	P
##	7	1	2	accuracy	binary	0.566	10	0.0344	P
##	8	1	2	roc_auc	binary	0.624	10	0.0414	P
##	9	2	2	accuracy	binary	0.590	10	0.0305	P
##	10	2	2	roc_auc	binary	0.612	10	0.0437	P

```
## # ... with 20 more rows
```

Plot metrics of different models



Plot metrics of different models



Show models with best metrics

```
## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean      n std_err .c
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <dbl>
## 1     1     5 accuracy binary    0.595    10  0.0345 Pr
## 2     2     2 accuracy binary    0.590    10  0.0305 Pr
## 3     1     4 accuracy binary    0.581    10  0.0354 Pr
## 4     3     2 accuracy binary    0.580    10  0.0333 Pr
## 5     1     1 accuracy binary    0.576    10  0.0329 Pr
```

Pick one model with the best metrics

```
best_forest <- forest_res %>%  
  select_best("accuracy")  
best_forest
```

```
## # A tibble: 1 x 3  
##   mtry min_n .config  
##   <int> <int> <chr>  
## 1     1     5 Preprocessor1_Model13
```

Finalizing best model

```
final_wf <-  
  tune_wf %>%  
  finalize_workflow(best_forest)
```

```
final_wf
```

```
## == Workflow =====  
## Preprocessor: Recipe  
## Model: rand_forest()  
##  
## -- Preprocessor -----  
## 2 Recipe Steps  
##  
## * step_rm()  
## * step_corr()  
##  
## -- Model -----  
## Random Forest Model Specification (classification)  
##
```

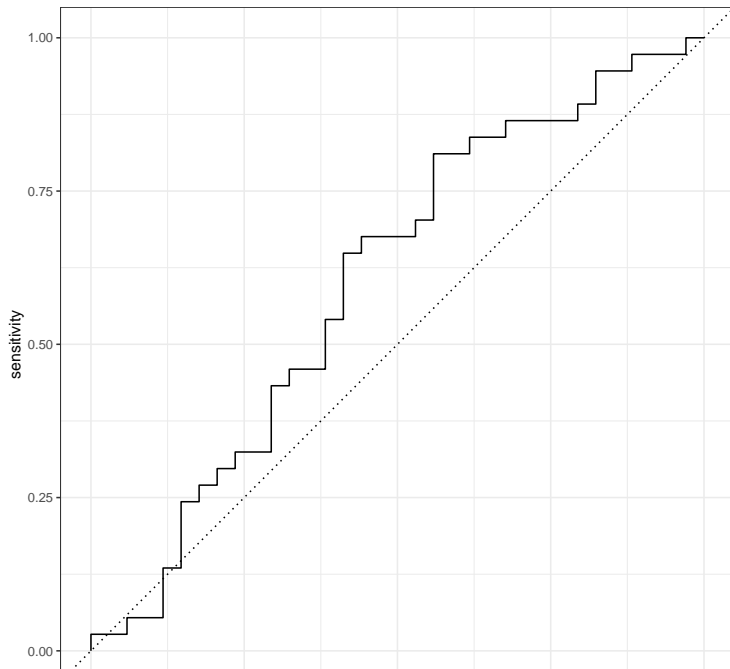
The final fit

```
final_fit <-  
  final_wf %>%  
  last_fit(vps_split)
```


```
final_fit %>%  
  collect_metrics()
```

```
## # A tibble: 2 x 4  
##   .metric .estimator .estimate .config  
##   <chr>    <chr>          <dbl> <chr>  
## 1 accuracy binary          0.606 Preprocessor1_Model1  
## 2 roc_auc  binary          0.603 Preprocessor1_Model1
```

The final fit ROC curve



The final fit Confusion Matrix



Model deployment

Model deployment

```
production <- read_csv("./data/vps_test_data.txt")

tmp <- final_wf %>% fit(vps_train) %>% predict(production) %>%
  rename(is_churn = .pred_class) %>%
  mutate(is_churn = ifelse(is_churn == "Yes", 1, 0 ))

production %>% select(! one_of('is_churn')) %>%
  bind_cols(tmp) %>%
  write_csv("./data/vps_test_data_pred_part_3.txt")
```

Further steps

Further steps

- ▶ demonstrate how to use **SparkR** (R on Spark)